



IT infrastructure refers to the composite hardware, software, network resources and services required for the existence, operation and management of an enterprise IT environment. It allows an organization to deliver IT solutions and services to its employees, partners and/or customers and is usually internal to an organization and deployed within owned facilities:



The underlying problem is on how to maintain the **STATE** of the servers in terms of what packages to be installed/removed, which services to started/stopped, creating user accounts, giving permissions, creating dir/files, taking backup etc.

## Configuration Management

Configuration management (CM) refers to the process of systematically handling changes to a system in a way that it maintains integrity over time

### What is Ansible?

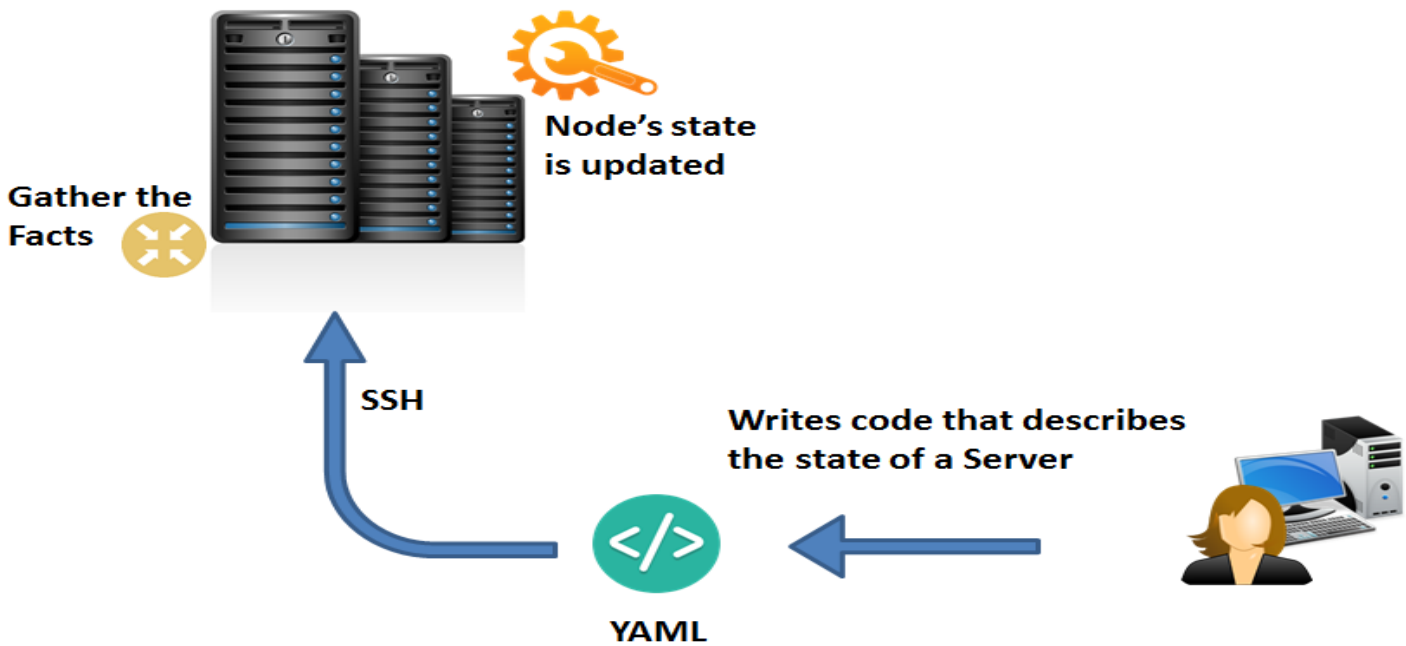
- Ansible is an automation engine that automates software provisioning, configuration management, and application deployment
- Manages infrastructure whether it is on-premises or in the cloud.
- It turns your infrastructure as code i.e your computing environment has some of the same attributes as your application:
  - Your infrastructure is versionable.
  - Your infrastructure is repeatable.
  - Your infrastructure is testable.
- You only need to tell what the desired configuration should be, not how to achieve it

### Why Ansible?

- Agentless
- Relies on ssh
- Uses python
- Push mechanism

### How Ansible Works?

Ansible works by connecting to your nodes and pushing out small programs, called "Ansible modules" to them. Ansible then executes these modules (over SSH by default), and removes them when finished. Your library of modules can reside on any machine, and there are no servers, daemons, or databases required.



The master node in the above picture is the controlling node (managing node) which controls the entire execution. It's the node from which you are running the installation. The inventory file provides the list of hosts where the Ansible modules need to be run and the management node does a SSH connection and executes the small modules on the hosts machine and installs the product/software.

Beauty of Ansible is that it removes the modules once those are installed so effectively it connects to host machine, executes the instructions and if it's successfully installed removes the code which was copied on the host machine which was executed

### Setup Ansible on AWS(Ubuntu Server)

```
$ apt-get update  
$ apt-get install -y ansible
```

### Setup Ansible on AWS(Centos Server)

```
$ yum install epel-release  
$ yum update  
$ yum install git python python-devel python-pip openssl ansible
```

## Ansible Environment Setup

- Settings in Ansible are adjustable via a configuration file called “ansible.cfg”. Edit `/etc/ansible/ansible.cfg` & enable the below lines

`inventory = /etc/ansible/hosts`

`sudo_user = root`

- Ansible recognizes systems listed in Ansible’s inventory file, which defaults to being saved in the location `/etc/ansible/hosts`
- The format for `/etc/ansible/hosts` is an INI-like format and looks like this:

`[groupname]`

`machinename | machineIP    #Incase of VM#`

`<MachineName> ansible_host=<<ec2-private-ip>> ansible_user=<<ec2-user>>  
ansible_ssh_private_key_file=/location/of/the/keypair/your-key.pem #Incase of  
AWS EC2#`

## Host Patterns

- Patterns in Ansible are how we decide which hosts to manage or what machines Ansible should connect
- A pattern can usually refer to a particular machine or an groupname
- "all" pattern refers to all the machines in an inventory
- You can refer to hosts within the group by adding a subscript to the group name while giving the pattern
  - `groupname[0]` -- picks the first machine in the group
  - `groupname[1]` -- picks the second machine in the group
  - `groupname[-1]` -- picks the last machine in the group
  - `groupname[0:1]` -- picks first 2 machine in the group

## Ansible Ad-Hoc Commands

- Ad hoc commands are commands which can be run individually to perform quick task only once.
- Use `/usr/bin/ansible` to run ad-hoc tasks really quick & don't want to save for later
- These are quick one-liner without writing a playbook

### Syntax:

**ansible [group|host|all] -m <module\_name> -a <arbitrary\_cmds>**

- List all the hosts under group 'demo'

**\$ ansible demo --list-hosts**

- Ping all the machines under group 'demo'

**\$ ansible demo -m ping**

- List all the files under `/home/ansible` dir on all the machines under group 'demo'

**\$ ansible demo -a "ls -al /home/ansible"**

- Display the last 10 lines from `/var/log/messages` file on all the machines under group 'demo'

**\$ ansible demo -a "cat /var/log/messages"**

- Run any tasks with sudo privilege, use `-s`

**\$ ansible demo -s -a "cat /var/log/messages"**

- Use copy module to copy a file from Ansible master to machines under group 'demo'

**\$ ansible demo -m copy -a "src=filename dest=filename"**

- Install/Remove/Update a Package

```
$ ansible all -s -m yum -a "pkg=httpd state=present"
```

```
$ ansible all -s -m yum -a "pkg=httpd state=absent"
```

```
$ ansible all -s -m yum -a "pkg=httpd state=latest"
```

- Start/Stop/Restart a Service

```
$ ansible all -s -m service -a "name=httpd state=started"
```

```
$ ansible all -s -m service -a "name=httpd state=stopped"
```

```
$ ansible all -s -m service -a "name=httpd state=restarted"
```

- Create/Delete a User account

```
$ ansible all -s -m user -a "name=adam"
```

```
$ ansible all -s -m user -a "name=adam state=absent"
```

### Gathering Facts:

- List all the properties that Ansible gathers while connecting any machine, output will be json format

```
$ ansible demo -m setup
```

## Ansible YAML Basics

- Ansible uses YAML syntax for expressing Ansible playbooks
- For Ansible, nearly every YAML file starts with a list
- Each item in the list is a list of key/value pairs, commonly called a "hash" or a "dictionary"
- All YAML files can optionally begin with "---" and end with "..."
- All members of a list are lines beginning at the same indentation level starting with a "- "

```
--- # A list of tasty fruits
```

```
fruits:
```

- Apple
- Orange
- Strawberry
- Mango

- A dictionary is represented in a simple key: value form (the colon must be followed by a space)

```
--- # An employee record
```

```
Employee:
```

```
  name: ADAM
```

```
  job: DevOps Engineer
```

```
  skill: Elite
```

## Ansible Playbooks

- Playbooks are Ansible's configuration, deployment, and orchestration language
- Playbooks describe a policy you want your remote systems to enforce, or a set of steps in a general IT process.
- Playbooks orchestrate steps of any manual ordered process, even as different steps must bounce back and forth between sets of machines in particular orders
- Playbooks are written in YAML format
- `/usr/bin/ansible-playbook` is used for running configurations from an playbook

### Syntax:

`Ansible-playbook <playbook>.yml`

Playbooks are divided sections & there are 3 major sections:

1. **Target Section** - Defines the hosts against which playbooks tasks has to be executed
2. **Variable Section** - Defines variables
3. **Tasks Section** - List of all modules that we need to run, in an order

### Target Section:

`---# My first Yaml`

`- hosts: <host_pattern>`

`become: <yes|no> # default is no #`

`become_user: <username> # user as whom ansible should be executed #`

`connection: <ssh|local> # defaults to ssh #`

`gather_facts: <yes|no> # defaults to yes #`



## Task Section:

```
--- # My First YAML playbook
```

```
- hosts: <group>
```

```
  become: <yes|no>
```

```
  connection: ssh
```

```
  gather_facts: no
```

```
  tasks:
```

```
    - name: <name of the task>
```

```
      <modulename>: <arbitrary commands>
```

## Variables Section:

- Refer various items for debug, set constant instead of typing every time
- foo\_port is a great variable. foo5 is fine too.
- foo-port, foo port, foo.port and 12 are not valid variable names.
- To use the variable, use the syntax '{{variablename}}'

```
--- # My First YAML playbook
```

```
- hosts: <group>
```

```
  become: <yes|no>
```

```
  connection: ssh
```

```
  gather_facts: no
```

```
  vars:
```

```
    <variablename>: <value>
```

```
  tasks:
```

```
    - name: <name of the task>
```

```
      <modulename>: <arbitrary commands>
```

## Handler Section:

- Consists the ability to notify a handler only when state change happens
- Also call another set of tasks

--- # My First YAML playbook

- hosts: <group>

become: <yes|no>

connection: ssh

gather\_facts: no

tasks:

- name: <name of the task>  
 <module name>: <arbitrary commands>  
 notify: <Handler task name>

handlers:

- name: <name of the handler task>  
 <module name>: <arbitrary commands>

## Ansible Dryrun

- Check whether the playbook is formatted correctly
- Test how the playbook is going to behave without running the tasks

\$ ansible-playbook playbook.yml --check

### Example:

```
--- # My First YAML playbook
```

```
- hosts: demo
```

```
  become: yes
```

```
  vars:
```

```
    pk: httpd
```

```
  tasks:
```

```
    - name: Install HTTPD server on centos 7
```

```
      yum: name='{{pk}}' state=installed
```

```
      notify: Restart HTTPD # this is called only if the action is ran & successful #
```

```
  handlers:
```

```
    - name: Restart HTTPD # this has to match the notify name #
```

```
      action: service name='{{pk}}' state=restarted
```

Run ansible-playbook to call the playbook

```
$ ansible-playbook playbook.yml
```

## Asynchronous Actions and Polling

- While using Ansible against multiple machines, the operations may run longer than SSH
- While one long task is running, another short task can be executed in asynchronous mode
- Specify the maximum runtime to timeout & how frequently to poll for status
  - ✓ **async**: <seconds to timeout the task>
  - ✓ **poll**: <seconds to poll for the status of the task>

### Syntax:

--- #Playbook to run task in parallel

- hosts: <group>

become: <yes|no>

connection: ssh

gather\_facts: no

tasks:

- name: <name of the task>

<modulename>: <arbitrary commands>

async: <seconds>

poll: <seconds>

**Example:** Check class notes for example

### Run Once

- In some cases there may be need to only run a task one time & on one host
- This can be achieved by configuring "run\_once" on a task
- This can be optionally paired with "delegate\_to" to specify an individual host to execute on

## Syntax:

```
--- #Playbook to run task once
- hosts: <group>
  become: <yes|no>
  tasks:
    - name: <name of the task>
      <modulename>: <arbitrary commands>
      run_once: true
      delegate_to: <node to which the task should run>
```

**Example:** Check class notes for example

## Loops

- Often you'll want to do many things in one task, such as create a lot of users, install a lot of packages, or repeat a polling step until a certain result is reached
- In those scenarios you will iterate the same task multiple times against different values using "with\_items" & read the value of each iteration using the ansible variable "item"

### Syntax:

```
--- #Playbook to run task in loop
- hosts: <group>
  become: <yes|no>
  tasks:
    - name: <name of the task>
      <modulename>: <arbitrary commands>
      with_items:
        - Value1
        - Value2
```

**Example:** Check class notes for example

## Conditions

- Few tasks might be needed to execute only on specific scenario
- Sometimes you will want to skip a particular step on a particular host
- In those cases we would use “**when**” statement

### Syntax:

```
--- #Playbook to run task based on a condition
- hosts: <group>
  become: <yes|no>
  tasks:
    - name: <name of the task>
      <modulename>: <arbitrary commands>
      when: <condition to satisfy on when to run the task>
```

**Example:** Check class notes for example

## Capture the task output

- By default Ansible do not capture the output of the tasks
- We have to explicitly store the output into a user variable using the “**register**” statement
- The values will be stored in json format

### Syntax:

```
--- #Playbook to run a task and capture its output
- hosts: <group>
```

become: <yes|no>

tasks:

- name: <name of the task>

    <modulename>: <arbitrary commands>

register: <variablename>

- debug: var=<variablename>.<attribute>

**Example:** Check class notes for example

## Error handling

- By default Ansible stops the execution of the playbook when it finds the first error, so if the first task fails to execute then it won't proceed further
- We can skip the failure by using “ignore\_errors” statement

### Syntax:

--- #Playbook to ignore errors on task1 and continue running task2

- hosts: <group>

become: <yes|no>

tasks:

- name: <name of the task1>

    <modulename>: <arbitrary commands>

    ignore\_errors: yes

- name: <name of the task2>

    <modulename>: <arbitrary commands>

**Example:** Check class notes for example

## wait\_for - Waits for a condition before continuing

- You can wait for a set amount of time to complete a task
- Waiting for a port to become available is useful for when services are not immediately available

### Syntax:

--- #Playbook to wait for a port to be available

- hosts: <group>

become: <yes|no>

tasks:

- name: <name of the task1>

<modulename>: <arbitrary commands>

- name: <name of the task to wait for a port available>

wait\_for:

port: <port#>

state: started

- name: <name of the task to wait for a file to be available>

wait\_for:

path: <filepath>

**Example:** Check class notes for example

### Tags

- If you have a large playbook it may become useful to be able to run a specific part of the configuration without running the whole playbook
- Use the statement “tags” to add a name to a task
- A task can have multiple tag names
- Same tag name can be share with multiple tasks to group them



## Syntax:

--- #Playbook to wait for a port to be available

- hosts: <group>

become: <yes|no>

tasks:

- name: <name of the task1>

<modulename>: <arbitrary commands>

tags:

- <tagname1>

- name: <name of the task1>

<modulename>: <arbitrary commands>

tags:

- <tagname2>

- If you want to run a playbook without certain tasks

**\$ ansible-playbook playbook.yml --tags "tagname"**

- If you want to run a playbook without certain tasks

**\$ ansible-playbook playbook.yml -skip-tags "tagname"**

**Example:** Check class notes for example

## Ansible Vault

- Ansible allows keeping sensitive data such as passwords or keys in encrypted files, rather than as plaintext in your playbooks
- To Run the playbook which is password protected using Ansible vault, use “--ask-vault-pass” while calling the playbook or “--vault-password-file FILE”. Where FILE is the name of file in which password is stored
- Creating a new Encrypted Files  
**\$ ansible-vault create playbook.yml**
- Edit the Encrypted File  
**\$ ansible-vault edit playbook.yml**
- Change the password  
**\$ ansible-vault rekey playbook.yml**
- Uncrypt the file  
**\$ ansible-vault decrypt playbook.yml**
- Encrypt an existing file  
**\$ ansible-vault encrypt playbook.yml**

## Ansible Roles

- Adding more & more functionality to the playbooks will make it difficult to maintain in a single file
- We can organize playbooks into a directory structure called roles
- This is already possible by ‘include’ directives however Roles are automation around it
- Default path for Roles **/home/ansible/playbooks/roles:/etc/ansible/roles:<PWD>**
- We can alternatively keep the master playbook in a different location & specify the Role path in ansible.cfg
- In the **/etc/ansible/ansible.cfg**, uncomment **roles\_path** & add the roles dir  
**roles\_path = /home/ansible/playbooks/roles**

## Framework/Syntax:

masterplaybook.yml

roles/<rolename>/

tasks/main.yml

vars/main.yml

handlers/main.yml

default/main.yml

meta/main.yml

## Syntax:

--- #Playbook for calling a role

- hosts: <group>

become: <yes|no>

roles:

- <Rolename>

Example: Check class notes for example

## Connecting to AWS using Ansible

- Connecting to AWS programmatically requires specific AWS Access Key Id & Secret Access Key
- Create ~/.boto & put the values obtained from the below steps:

### [Credentials]

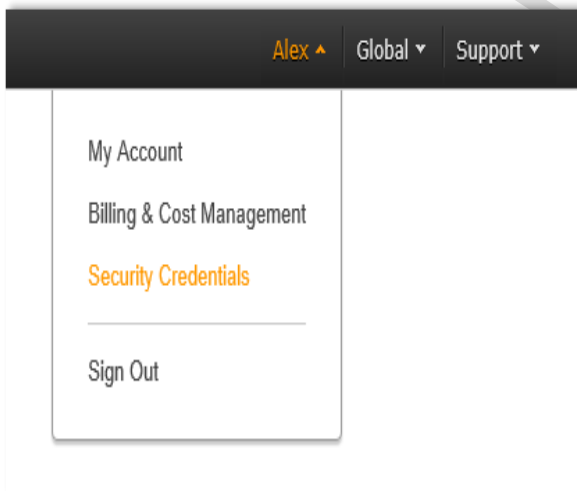
aws\_access\_key\_id = <your\_access\_key\_here>

aws\_secret\_access\_key = <your\_secret\_key\_here>

- Copy the keypair.pem file to the Ansible Master server

## How to find your AWS Access Key ID and Secret Access Key

- ✓ Go to Amazon Web Services console and click on the name of your account (it is located in the top right corner of the console). Then, in the expanded drop-down list, select Security Credentials.



- ✓ Click the Continue to Security Credentials button

You are accessing the security credentials page for your AWS account. The account credentials provide unlimited access to your AWS resources.

To help secure your account, follow an [AWS best practice](#) by creating and using AWS Identity and Access Management (IAM) users with limited permissions.

[Continue to Security Credentials](#) [Get Started with IAM Users](#)

☐ Don't show me this message again

- ✓ Expand the Access Keys (Access Key ID and Secret Access Key) option. You will see the list of your active and deleted access keys

## Your Security Credentials

Use this page to manage the credentials for your AWS account. To manage credentials for AWS Identity and Access Management (IAM) users, use the [IAM Console](#).  
To learn more about the types of AWS credentials and how they're used, see [AWS Security Credentials](#) in AWS General Reference.

- + Password
- + Multi-Factor Authentication (MFA)
- Access Keys (Access Key ID and Secret Access Key)

You use access keys to sign programmatic requests to AWS services. To learn how to sign requests using your access keys, see the [signing documentation](#). For your protection, store your access keys securely and do not share them. In addition, AWS recommends that you rotate your access keys every 90 days.

Note: You can have a maximum of two access keys (active or inactive) at a time.

Created	Deleted	Access Key ID	Last Used	Last Used Region	Last Used Service	Status	Actions
Jun 18th 2015		AKIAIUZCFUBHKB3TNOA	N/A	N/A	N/A	Active	<a href="#">Make Inactive</a>   <a href="#">Delete</a>
Jun 18th 2015	Jun 18th 2015	AKIAJKOPHEMD52BJRV6Q	N/A	N/A	N/A	Deleted	
Jun 18th 2015	Jun 18th 2015	AKIAJLWJENP4LYYG7U5Q	N/A	N/A	N/A	Deleted	
Jun 4th 2015	Jun 18th 2015	AKIAIV75YZYWGRZY33RA	N/A	N/A	N/A	Deleted	
Jun 15th 2015	Jun 18th 2015	AKIAIAWO43ZQO2PWNUIJQ	N/A	N/A	N/A	Deleted	
May 29th 2015	Jun 15th 2015	AKIAI63LOZJPNTZUIBA	N/A	N/A	N/A	Deleted	

[Create New Access Key](#)



### Important Change - Managing Your AWS Secret Access Keys

As described in a [previous announcement](#), you cannot retrieve the existing secret access keys for your AWS root account, though you can still create a new root access key at any time. As a [best practice](#), we recommend [creating an IAM user](#) that has access keys rather than relying on root access keys.

- + CloudFront Key Pairs
- + X.509 Certificates
- + Account Identifiers

- ✓ To generate new access keys, click the Create New Access Key button.

You use access keys to sign programmatic requests to AWS services. To learn how to sign requests using your access keys, see [Access Keys](#). You must rotate your access keys every 90 days.

Note: You can have a maximum of two access keys (active or inactive) at a time.

Created	Deleted	Access Key ID
Jun 24th 2015		AKIAJANMQRD73LQYU34A
Jun 18th 2015	Jun 24th 2015	AKIAIIUZCFUBHKB3TNOA
Jun 18th 2015	Jun 18th 2015	AKIAJKQPHEMD52BJRV6Q
Jun 18th 2015	Jun 18th 2015	AKIAJLWJENP4LYYG7U5Q
Jun 4th 2015	Jun 18th 2015	AKIAIV75YZYWGRZY33RA
Jun 15th 2015	Jun 18th 2015	AKIAIAWO43ZQO2PWNUJQ
May 29th 2015	Jun 15th 2015	AKIAI633LOZJPNTZUIBA

Create New Access Key



### Important Change - Managing Your AWS Secret Access Keys

As described in a [previous announcement](#), you cannot retrieve the existing secret access keys for your IAM user. You must create new access keys rather than relying on root access keys.

- ✓ Click Show Access Key to have it displayed on the screen. Note, that you can download it to your machine as a file and open it whenever needed. To download it, just click the Download Key File button.

### Create Access Key

☒ Your access key (access key ID and secret access key) has been created successfully.

Download your key file now, which contains your new access key ID and secret access key. If you do not download the key file now, you will not be able to retrieve your secret access key again.

To help protect your security, store your secret access key securely and do not share it.

▼ Hide Access Key

Access Key ID: AKIAI633LOZJPNTZUIBA

Secret Access Key:

Download Key FileClose

## Creating EC2 Instances in AWS using Ansible

```
--- # Creating EC2 Instances in AWS
- hosts: demo
  become: yes
  connection: local
  tasks:
    - name: Install python-pip library # prerequisite
      apt: name={{item}}
      with_items:
        - python-pip
        - python-dev
    - name: Install python-boto library
      pip: name=boto
    - name: Create AWS Instances
      ec2:
        key_name: "wezva"
        instance_type: "t2.micro"
        image: "ami-c58c1dd3"
        wait: true
        region: "us-east-1"
```

## Creating/Deleting S3 Buckets in AWS using Ansible

--- # Creating EC2 Instances in AWS

- hosts: demo

become: yes

connection: local

tasks:

- name: Install python-pip library # prerequisite

apt: name='{{item}}'

with\_items:

- python-pip

- python-dev

- name: Install python-boto library

pip: name=boto

- name: Create S3 Bucket

S3\_bucket:

name: mys3bucketfortest

region: "us-east-1"

state: present #use state as absent to delete the S3 bucket #

[www.wezva.com](http://www.wezva.com)

facebook

<https://www.facebook.com/wezva>

Linked in

<https://www.linkedin.com/in/wezva>



+91-9739110917

+91-9886328782